

RE030- Evaluation Metrics for Personalized Food Recommendation Systems

Submitted by – Prajjwal Makkar (Research Lead and Project Co-Lead)

Table of Contents

- 1. Introduction**
- 2. Types of Evaluation Metrics**
- 3. Comparison Across Applications**
- 4. Limitations of Evaluation Metrics**
- 5. Key Insights**
- 6. Implications for Food Remedy API**

Introduction

A crucial element of contemporary nutrition and health applications are personalized meal suggestion systems. These systems seek to facilitate better nutritional choices by analyzing user preferences, objectives, and behavioral patterns. However, in order to guarantee accuracy and relevance, their efficacy must be assessed using the proper evaluation techniques.

There are two primary reasons why evaluation is crucial. It makes sure suggestions meet user needs and assesses if they result in significant behavioral changes, like better adherence to dietary objectives and healthier choices. This study looks at evaluation metrics, compares how they are used in different applications, and investigates how they might be used to improve the Food Remedy API.

Types of Evaluation Metrics

a) Accuracy Metrics

Metrics for accuracy assess how well the recommendation system anticipates items that are pertinent to a user. Precision, recall, and relevance scoring are common metrics. Recall gauges how well the system captures all pertinent possibilities, whereas precision counts the percentage of suggested items that are truly helpful.

Accuracy in food applications involves not just matching tastes but also adhering to dietary constraints, such as nutritional balance, calorie objectives, or allergies.

But accuracy is not enough on its own. Technically sound advice may be made by a machine, but it may not engage consumers or change their behavior. As a result,

accuracy should be viewed as a fundamental statistic rather than a comprehensive indicator of success.

b) Engagement Metrics

The way users engage with the recommendation system is measured by engagement metrics. These consist of interaction frequency, time spent viewing suggestions, and click-through rate (CTR).

Recommendations that are well-presented, pertinent, and aesthetically pleasing are usually associated with high engagement.


But meaningful results are not necessarily the result of interaction. Therefore, in order to fully comprehend their influence, engagement indicators should be read in conjunction with behavioral data.

c) Behavior and outcome metrics

The practical effects of recommendations are the main focus of behavioral metrics. These include long-term behavioral changes, adherence to dietary regimens, and conversion rates (e.g., choosing a suggested item).

These statistics assess whether users follow recommendations and make better decisions, in contrast to engagement measures.

Because they directly reflect the goal of the system, which is to improve user outcomes, behavioral metrics are the most valuable. But because they necessitate longer-term monitoring and more in-depth study, they are also the hardest to quantify.



Comparison Across Applications

Depending on their primary goals, target audience, and philosophy of product design, various applications prioritize different evaluation measures. Gaining knowledge about how these platforms assess efficacy can help one better understand how personalized recommendation systems function in practical situations.


1) MyFitnessPal

MyFitnessPal primarily focuses on engagement and behavioural tracking metrics. Its recommendation system is closely tied to food logging and daily calorie tracking.

The platform evaluates success based on:

- a) Frequency of user interactions (daily log-ins, food entries)
- b) Consistency in tracking behaviour over time
- c) User retention and long-term engagement

The system does not rely heavily on traditional recommendation accuracy metrics. Instead, it measures effectiveness through habit formation and



continued usage.

2) Yuka

Yuka places a lot of emphasis on measures based on trust, accuracy, and transparency. Its system uses a transparent scoring system to assess food products according to their overall health impact, additives, and nutritional quality.

Success is determined by:

- a) Precision of product assessment
- b) Scoring and explanation clarity
- c) User confidence in the data supplied

Yuka places more focus on decision support than interaction frequency.

3) **Fastic**

Fastic emphasizes outcome-based and behavioral measurements, especially when it comes to fasting and developing habits.

Its assessment methodology consists of:

- a) Observance of fasting regimens
- b) Rates of fasting cycle completion
- c) Long-term constancy of habits

Instead of only making suggestions, the system is intended to affect behavior.

Comparative analysis

There are noticeable variations in the prioritization of evaluation metrics across these applications:

- a) MyFitnessPal places a strong emphasis on participation and regularity
- b) Yuka places a high value on precision and openness
- c) Fastic emphasizes the development of habits and behavioral consequences

This comparison demonstrates that recommendation systems cannot be adequately assessed using a single metric.

One important finding is that:

- a) Engagement measures are simpler to gauge
- b) Accuracy measurements increase trust
- c) Behavioral indicators provide the most insightful data

Key Takeaways from Comparison

- a) Instead of depending just on one assessment method, effective systems combine several.
- b) Short-term engagement should not take precedence over behavioral outcomes.
- c) Transparency increases trust, but it needs to be combined with practical suggestions.
- d) To determine actual impact, long-term tracking is necessary.

Relevance to Food Remedy API

This comparison indicates that a hybrid evaluation strategy is necessary for the Food Remedy system.

The system ought to:


- a) Make sure recommendations are pertinent by using accuracy metrics

- b) Utilize engagement metrics to gauge usability and interactivity
- c) Track real-world impact with behavioral metrics

Limitation of Evaluation Metrics

Although evaluation measures are crucial for evaluating recommendation systems, they have a number of drawbacks. As users may interact with recommendations without following them, metrics like click-through rate and engagement frequency may not necessarily indicate relevant consequences. In a similar vein, accuracy measurements might show that suggestions are pertinent, but they don't ensure that users would implement them in practical situations.

The fact that different users react to suggestions in different ways is another significant drawback. It is challenging to generalize outcomes from a system that works well for one group but not for another. Furthermore, even if they are the most beneficial, behavioral measurements are more difficult to evaluate precisely and need long-term observation. These difficulties demonstrate the necessity of a well-rounded assessment strategy that incorporates a variety of measures to more accurately gauge system efficacy.



Key Insights-

A single statistic is insufficient to evaluate the efficacy of personalized meal recommendation systems, as the evaluation metrics study reveals. Although accuracy measurements guarantee the relevance of recommendations, they do not ensure user adoption. Although engagement metrics provide light on user involvement, they don't indicate whether or not users make better choices. Because they evaluate real-world

results, behavioral metrics provide the most meaningful evaluation; yet, they are more difficult and necessitate long-term surveillance.

A crucial realization is that, rather than depending just on one method, efficient evaluation necessitates a combination of various criteria. To fully comprehend performance, systems must strike a balance between accuracy, engagement, and behavioral results.

Furthermore, since the ultimate objective of recommendation systems is to encourage persistent, healthy eating choices, long-term user behavior is more significant than short-term involvement.

Implications for Food Remedy

The analysis's conclusions recommend that the Food Remedy API use a hybrid evaluation strategy that incorporates behavioral, accuracy, and engagement criteria. While engagement metrics can be used to assess how users connect with the system, accuracy measurements can guarantee that recommendations match user preferences and dietary needs. However, behavioral measurements should be the main focus since they show if the system is effective in influencing actual food choices.

Evaluation for Food Remedy should go beyond documenting interactions. Actions like choosing suggested foods, switching to healthier foods frequently, and consistently adhering to dietary recommendations should be tracked by the system. Monitoring these behaviors over time will shed more light on how well the recommendation system works.

Continuous improvement can also be made possible by using analytics and feedback loops. The system can improve its scoring and suggestion systems to produce more pertinent and significant outcomes by examining user reactions to recommendations. This strategy guarantees that the system leads to significant improvements in user health outcomes in addition to performing well technically.

Appendix A — Key Engagement and Evaluation Metrics

Metric	Definition / Formula	Example / Interpretation
Click-Through Rate (CTR)	$\text{CTR} = \frac{\text{Clicks}}{\text{Impressions}}$	50 clicks / 1000 impressions = 0.05 (5 %). Higher CTR → more relevant content.
View-Through Rate (VTR)	$\text{VTR} = \frac{\text{Completed Views}}{\text{Impressions}}$	2150 completed views / 18 700 impressions ≈ 11.5 %. Shows content completion.
Normalized Discounted Cumulative Gain (NDCG@k)	$\text{NDCG@k} = \frac{\sum_{i=1}^k \frac{2^{rel_i} - 1}{\log_2(i+1)}}{\text{IDCG@k}}$	Ranges 0–1. 1.0 = perfect ranking; accounts for position and relevance grades.
Dwell Time	$t_{exit} - t_{entry}$	If user opens 12:00:00 → leaves 12:00:45 → 45 s dwell. Indicates engagement depth.
Completion Rate	$\text{Completion Rate} = \frac{\text{Video Completes}}{\text{Video Starts}}$	45 completed / 100 starts = 45 %. Higher means viewers finish content.
Expected Calibration Error (ECE)	$\text{ECE} = \sum_{m=1}^M \frac{n_m}{N} p_m - o_m $	o_m - p_m
Area Under the Uplift Curve (AUUC)	$\text{AUUC} = \int_0^1 U(x) dx$	Measures total incremental gain achieved by targeting top-uplift users. Higher AUUC = better uplift ranking.